

Performance Considerations on NUMA Machines

Basics of Parallel Computing

Assoc.Prof. Dr. Sascha Hunold


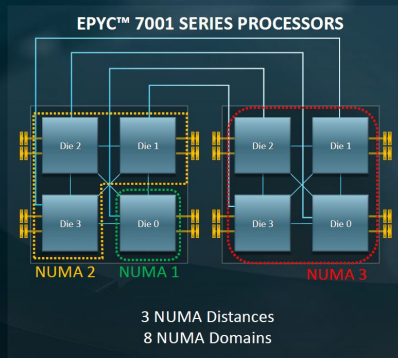
TU Wien



Informatics

NUMA Performance in Practice

EPYC™ 7002 SERIES NUMA ADVANCEMENTS



Reduced Number of NUMA Domains and Distances: Improved NUMA Attributes for General Workloads

source: <https://www.servethehome.com/wp-content/uploads/2019/08/AMD-EPYC-7002-Architecture-NUMA-Reduction-to-104ns-Close-201ns-Far.jpg>

- our machine **nebula**


```
hunold@nebulac:~$ lscpu |grep "Model name"  
Model name:          AMD EPYC 7551 32-Core Processor
```

- <https://www.amd.com/de/products/cpu/amd-epyc-7551>
 - Product line: AMD EPYC 7001 Series
- `hwloc` can help to reveal processor details and topology


```
hunold@nebulac:~$ hwloc-ls - --of pdf > nebulac_node.pdf
```

NUMA Examples / AMD EPYC / Intel Xeon III

- one NUMA node in detail
- notice the core numbering!
 - cores 0, 16, 32, 48 share one L3 cache (important)



in total: 8 NUMA nodes (2×4)



Experiment

- we check the bandwidth (how many bytes we can read and write from/to memory)
 - we say bandwidth but actually mean throughput
 - bandwidth is theoretical, throughput is a practical measure
- we use 2 threads and **nebula**

config 1

- we pin both threads to **one NUMA** node

config 2

- we pin both threads to **different NUMA** nodes

hypothesis

- config 2 gives more bandwidth as we can leverage more memory controllers

config 1 (same NUMA node)

```
hunold@nebulac:~/tmp/stream$ numactl --physcpubind=0,16 ./stream_omp
```

```
-----  
Number of Threads requested = 2
```

```
Number of Threads counted = 2
```

```
-----  
Function      Best Rate MB/s  Avg time     Min time     Max time  
Copy:         20735.7    0.007731    0.007716    0.007753  
Scale:        20562.8    0.007788    0.007781    0.007798  
Add:          22720.5    0.010576    0.010563    0.010584  
Triad:        22735.9    0.010566    0.010556    0.010578
```

- memory throughput: ~23 GB/s

config 2 (different NUMA nodes)

```
hunold@nebulac:~/tmp/stream$ numactl --physcpubind=0,2 ./stream_omp
```

```
-----  
Number of Threads requested = 2
```

```
Number of Threads counted = 2
```

```
-----  
Function      Best Rate MB/s  Avg time     Min time     Max time  
Copy:         30805.1    0.005221    0.005194    0.005262  
Scale:        31658.1    0.005064    0.005054    0.005077  
Add:          39435.6    0.006102    0.006086    0.006113  
Triad:        38609.7    0.006235    0.006216    0.006257
```

- memory throughput: ~39 GB/s

■ memory latencies between NUMA nodes

```
hunold@nebulac:~$ numactl -H
available: 8 nodes (0-7)
node 0 cpus: 0 8 16 24 32 40 48 56
node 0 size: 31934 MB
node 0 free: 6940 MB
node 1 cpus: 2 10 18 26 34 42 50 58
node 1 size: 32252 MB
node 1 free: 30238 MB
node 2 cpus: 4 12 20 28 36 44 52 60
node 2 size: 32252 MB
node 2 free: 31221 MB
// .. I removed a few lines
node distances:
node   0   1   2   3   4   5   6   7
 0:  10  16  16  16  28  28  22  28
 1:  16  10  16  16  28  28  28  22
 2:  16  16  10  16  22  28  28  28
 3:  16  16  16  10  28  22  28  28
 4:  28  28  22  28  10  16  16  16
 5:  28  28  28  22  16  10  16  16
 6:  22  28  28  28  16  16  10  16
 7:  28  22  28  28  16  16  16  10
```

- let us look at a **hydra** node

```
hunold@hydra01:~$ lscpu | grep "Model name"
Model name:          Intel(R) Xeon(R) Gold 6130F CPU @ 2.10GHz

hunold@hydra01:~$ hwloc-ls - --of pdf > hydra_node.pdf
```

NUMA Examples / AMD EPYC / Intel Xeon IX



■ thus, sockets are NUMA nodes

- memory latency to another NUMA node is roughly twice higher

```
hunold@hydra01:~$ numactl -H
available: 2 nodes (0-1)
node 0 cpus: 0 2 4 6 8 10 12 14 16 18 20 22 24 26 28 30
node 0 size: 46803 MB
node 0 free: 42308 MB
node 1 cpus: 1 3 5 7 9 11 13 15 17 19 21 23 25 27 29 31
node 1 size: 48360 MB
node 1 free: 46808 MB
node distances:
node    0    1
  0:   10   21
  1:   21   10
```

Does this inter-core latency matter? I

Experiment

- OSU MPI Microbenchmarks
- 2 processes perform a **ping-pong**
 - A sends message to B, B sends message back
 - we measure the (mean) time to complete this ping-pong
- we do this on one compute node of **hydra**

config 1

- processes 1 and 2 are mapped to **same socket**

config 2

- processes 1 and 2 are mapped to **different sockets**

hypothesis

- config 2 is slower than config 1 (at least for small messages)
- since physical latency within one socket should be smaller (see slide before)

Does this inter-core latency matter? II

■ config 1 (on same socket)

```
srunk --nodelist=hydra01 --ntasks-per-node=2 \  
-m block:block ./mpi/pt2pt/osu_latency -m 1:128 \  
-i 10000
```

OSU MPI Latency Test v5.4.1

# Size	Latency (us)
1	0.41
2	0.41
4	0.40
8	0.41
16	0.41
32	0.56
64	0.55
128	0.58

■ indeed, config 1 is faster

■ take away message: **placement** of processes and threads (on which core) is **important**

■ config 2 (on different sockets)

```
srunk --nodelist=hydra01 --ntasks-per-node=2 \  
-m block:cyclic ./mpi/pt2pt/osu_latency -m 1:128 \  
-i 10000
```

OSU MPI Latency Test v5.4.1

# Size	Latency (us)
1	0.65
2	0.65
4	0.65
8	0.65
16	0.65
32	1.01
64	1.00
128	1.01