

Scalability, Cost, and Work

Basics of Parallel Computing

Assoc.Prof. Dr. Sascha Hunold

TU Wien



Informatics

Gustafson-Barsis's Law

Gustafson-Barsis's Law

- Gustafson, 1988, re-evaluation of Amdahl's law
- observation of Gustafson's Law: **more powerful computer systems** usually **solve larger problems**, not the same size problem in less time

Gustafson's argument is

- a program consists of **serial fraction s** and a **parallel fraction $1 - s$**
- the **parallel portion scales perfectly** with the number of processors p
- if the parallel part is run on a serial processor, it would take p times longer than on p processors
- the speed-up is then the following (recall that $s + (1 - s) = 1$)

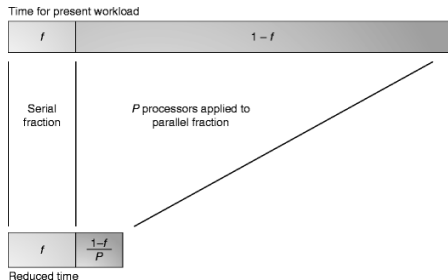
$$S = \frac{s + p(1 - s)}{s + (1 - s)} = s + p(1 - s)$$

- shows that the serial fraction does not theoretically limit the speed-up, if the problem (workload) scales with the number of processors

Amdahl vs. Gustafson

Amdahl

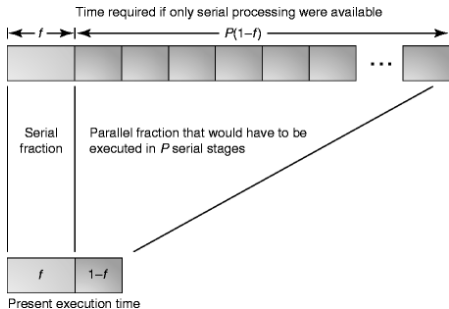
- Amdahl's Law **fixes the problem size**
- How much faster will my program run with a fixed problem size if I use p processors?



source: Encyclopedia of Parallel Computing


Gustafson

- Gustafson's Law **fixes the run time**
- How much longer does it take for a given workload to be executed if run sequentially?



Gutafson's Law Visualized

```
1 library(ggplot2)
2 su <- function(p, seq_share) {
3   seq_share + p * (1-seq_share)
4 }
5
6 df <- expand.grid(p = c(1,2,8,16,128,256,1024), seq = c(0.5,0.25,0.1,0.01))
7 df$su <- su(df$p, df$seq)
8
9 ggplot(df, aes(x=p, y=su, color=factor(seq))) + geom_line() +
10   xlab("number of processors") + ylab("speed-up") + theme_bw() + labs(color="s")
```



Strong and Weak Scaling

Strong and Weak Scaling

- in parallel computing, we typically use two analyses (depending on the type of problem)

Strong Scaling (based on Amdahl's law)

- keep **input size fixed**
- **increase** the **number of processors**

Weak Scaling (based on Gustafson's law)

- **increase** the **number of processors**
- **scale instance size** with the number of processors

Strong and Weak Scaling Example I

- we perform a **weak scaling** and **strong scaling** experiment with the same code
- we use our example from the previous lecture, the `dummy_count` in R

```
1 dummy_count <- function(max) {  
2   sum = 0  
3   for(i in 1:max) {  
4     sum = sum + i  
5   }  
6   sum  
7 }
```

Strong and Weak Scaling Example II

- in the **strong-scaling** experiment, we keep the **problem size fixed**
- thus, we fix `input <- 1:10000`
- thus, the first ten elements of `input` are
 - 1 2 3 4 5 6 7 8 9 10
- then, we increase the number of processors (cores)

```
1 for (i in p) {  
2   a <- within(benchmark(replications=c(10),  
3                       mclapply(input, dummy_count, mc.cores = i),  
4                       columns=c('test', 'elapsed', 'replications')),  
5               { average = elapsed/replications })  
6   df[df$p==i, ]$tstrong <- a$average[[1]]  
7 }
```

- we can then compute the **relative speed-up**
 - as usual $T_{\text{par}}(n, 1)/T_{\text{par}}(n, p)$
 - `df$sustrong <- df[df$p==1,]$tstrong / df$tstrong`

Strong and Weak Scaling Example III

- in the **weak-scaling** experiment, we keep the **problem size per core fixed**
 - thus, we grow the problem with an increasing number of cores
- with 1 core, we use `input <- 1:10000`
 - with 2 cores, we double the work to `c(input, input)`
 - with 3 cores, we triple the work to `c(input, input, input)`
- the experiment then looks like this

```
1 for (i in p) {  
2   a <- within(benchmark(replications=c(10),  
3                       mclapply(rep(input,i), dummy_count, mc.cores = i),  
4                               columns=c('test', 'elapsed', 'replications')),  
5         { average = elapsed/replications })  
6   df[df$p==i, ]$tweak <- a$average[[1]]  
7 }
```

- we compute a **scaled speed-up**, where we relate the time to finish the scaled workload on only 1 processor to the time it took to complete the scaled workload on p processors
- $$S = \frac{pT_{\text{par}}(n,1)}{T_{\text{par}}(pn,p)}$$
- `df$suscaled <- (df[df$p==1,]$tweak * df$p) / df$tweak`

Strong and Weak Scaling Example IV

hydra35


- we now perform this experiment on one compute node of the hydra cluster
- we select hydra35 (no particular reason, but we had to select one node)
- the experimental results are as follows

p	tstrong	tweak	sustrong	suscaled
1	1.57	1.57	1	1
2	0.79	1.59	1.98	1.98
4	0.4	1.59	3.91	3.94
16	0.12	1.63	13.18	15.38
32	0.1	1.7	16.24	29.54


Strong and Weak Scaling Example V

Strong Scaling - hydra35

```
1 ggplot(df, aes(x=p, y=tstrong)) +  
2   geom_point() + geom_line() +  
3   ylim(0,1.7) +  
4   ylab("time [s]")
```




```
1 ggplot(df, aes(x=p, y=sustrong)) +  
2   geom_point() + geom_line() +  
3   geom_line(data=df, aes(x=p, y=p), color="red") +  
4   ylab("speed-up")
```




Strong and Weak Scaling Example VI

Weak Scaling - hydra35

```
1 ggplot(df, aes(x=p, y=tweak)) +  
2   geom_point() + geom_line() +  
3   ylim(0,1.7) +  
4   ylab("time [s]")
```



```
1 ggplot(df, aes(x=p, y=suscaled)) +  
2   geom_point() + geom_line() +  
3   ylab("scaled speed-up")
```



A Note on Weak Scaling I

- we usually check whether the runtime stays constant when increasing both the problem size and the number of processors
- but we need to scale the problem size correctly
- in our previous example
 - the number of operations depends on the values in the list
 - for a list `1:n` the total number of operations is $O = \sum_{j=1}^n j$
 - there are j additions per invocation of `dummy_count`
 - with p processors, we simply multiply with p , i.e., pO

A Note on Weak Scaling II

- but what if we do a **matrix-vector product**?
 - **complexity** of a matrix vector product $O(n^2)$, for a matrix of size $n \times n$
- assume that we **parallelize the matrix-vector product**
 - each processor gets n/p rows of matrix A
 - each processor needs to perform $2n^2/p$ operations (ADD + MUL per element), which is its **work**
- if we want to **increase n with p** , we solve by n
 - $w_1 = w_p = \frac{2n^2}{p}$, $n = \sqrt{\frac{w_1 p}{2}}$
 - w_1 denotes the work done on one processor for the sequential run
 - w_p denotes the work done per processor in the parallel run
- now, we want to perform a **weak-scaling analysis** of our algorithm
 - we set $n = 100$, then $w_1 = 2n^2 = 20\,000$

```
1 n <- 100
2 w <- 2*n^2
3 df <- data.frame(p=c(1,2,4,8,16,32,64,128))
4 df$n <- ceiling(sqrt(w*df$p / 2))
5 print(toOrg(df))
```

p	n
1	100
2	142
4	200
8	283
16	400
32	566
64	800
128	1132

Embarrassing Parallelism


- We often hear the word “**embarrassingly parallel**”, but what does that mean?
- it means that the parallelization **strategy is obvious**
 - no advanced algorithms needed
 - very **little communication** overhead
 - little or **no effort to split up** work
- for example: parameter sweeps
 - program A is called with parameters p_1 and p_2
 - p_1 may have 30 different values
 - p_2 may have 20 different values
 - we got 600 different combinations of p_1 and p_2
 - we can parallelize easily over these combinations
 - Note: we could also parallelize A itself

Cost and Work of Parallel Programs

Cost and Work of Parallel Programs I

Cost

The **cost** of a parallel algorithm denotes the total time that one uses all p processors, i.e.,
 $C = pT_{\text{par}}(p, n)$.



Cost Optimality [1]

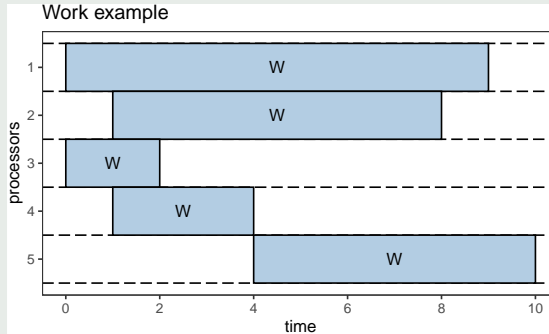
A parallel program is **cost optimal** if the cost of solving a problem on a parallel computer has the same asymptotic growth (as a function of the input size) as the fastest-known sequential algorithm on a single processing element, i.e.,

$$pT_{\text{par}}(p, n) \in O(T_{\text{seq}}(n)) \quad .$$

Cost and Work of Parallel Programs III

Work

The **work** W of a parallel algorithm denotes the **number of operations** that are carried out.




Work Optimality

A parallel program is **work optimal** if the work W has the same asymptotic growth as the fastest-known sequential algorithm on a single processing element, i.e.,


$$W \in O(T_{\text{seq}}(n)) \quad .$$

Cost and Work of Parallel Programs V


sequential



work optimal but not cost optimal



work and cost optimal



Cost Analysis / Example

Example of Cost Analysis - Attempt 1 I


- we consider the problem of adding n numbers in parallel
- we assume to have n processing elements (yes, as many processing elements as numbers) and we also assume that n is a power of two
- we can solve the problem in $\log n$ steps
 - propagate partial sums up a logical binary tree




(a) Initial data distribution and the first communication step



(b) Second communication step



(c) Third communication step



(d) Fourth communication step



(e) Accumulation of the sum at processing element 0 after the final communication

source: [1]

Example of Cost Analysis - Attempt 1 II

- sequential running time $\Theta(n)$
- parallel running time $\Theta(\log n)$
- speed-up: $\Theta\left(\frac{n}{\log n}\right)$

```
1 n <- c(2,4,16,128,1024, 1024^2)
2 df <- data.frame(n=n, su= n / log2(n))
3 df
```

	n	su
1	2	2.00000
2	4	2.00000
3	16	4.00000
4	128	18.28571
5	1024	102.40000
6	1048576	52428.80000

Example of Cost Analysis - Attempt 1 III

What about the cost?

- we use n processors for time $\Theta(\log n)$
- thus, the cost of the parallel algorithm is $\Theta(n \log n)$
- recall that the sequential running time is $\Theta(n)$
- we ask if $O(n \log n) \in O(n)$?
 - no, and thus, the algorithm is **not** cost-optimal

Attempt 2 - Scaling I

Can we improve this algorithm?

- using n processors for n elements unrealistic
 - although good to design algorithm
 - we can always **scale down** a parallel system: some processors simulate the work of other processors
- let us assume now that $p < n$
 - n and p are powers of two
- we use the same algorithm as before, but we **simulate** the previous n processors **on** p processors
- our previous processor i is now mapped to processor $i \bmod p$
 - e.g. $n = 16, p = 4$





```
1 df <- data.frame(i=seq(0,15))
2 df$p_real <- df$i %% 4
3 df
```

	i	p_real
1	0	0
2	1	1
3	2	2
4	3	3
5	4	0
6	5	1
7	6	2
8	7	3
9	8	0
10	9	1
11	10	2
12	11	3
13	12	0
14	13	1
15	14	2
16	15	3

Attempt 2 - Scaling II





- 1 virtual processor i is simulated by physical processor $i \bmod p$
- 2 in the first $\log p$ steps of the original algorithm on n processors now take $(n/p) \log p$ steps on p processors
- 3 in the remaining steps, no communication required (data already at physical processor), data added locally
- 4 our example, with $n = 16$ and $p = 4$

Attempt 2 - Scaling (Visualization)




- simulating 16 original processors on 4 processors
- first communication step
- original algorithm
 - $p1 \rightarrow p0$
 - $p3 \rightarrow p2$
 - $p5 \rightarrow p4$
 - $p7 \rightarrow p6$
 - ...
- now, $p0$ simulates $i \bmod 4$
 - it simulates $p0, p4, p8, p12$
 - thus, it needs 4 substeps
- similarly for the other processors

Attempt 2 - Scaling (Visualization)




- simulating 16 original processors on 4 processors
- second communication step
- original algorithm
 - $p_2 \rightarrow p_0$
 - $p_4 \rightarrow p_2$
 - $p_6 \rightarrow p_4$
 - $p_8 \rightarrow p_6$
 - ...
- now, p_0 simulates p_0, p_4, p_8, p_{12}
 - p_2 simulates p_2, p_6, p_{10}, p_{14}
 - thus, we need 4 substeps

Attempt 2 - Scaling (Visualization)




- subresults are now locally available on p0
- no further communication required

Attempt 2 - Scaling (Visualization)



- subresults are now locally available on p0
- no further communication required

Attempt 2 - Scaling (Visualization)



- final result
- runtime of original algorithm:
 $\log n = \log 16 = 4$
 - cost: $n \log n$
- runtime of simulation:
 - communication:
 $n/p \log p = 16/4 \log 4 = 4 \cdot 2 = 8$
 - computation: $n/p = 16/4 = 4$
 - overall: $O(n/p \log p + n/p) = O(n/p \log p)$
 - cost: $O(p \frac{n}{p} \log p) = O(n \log p)$
- thus, this strategy is also **not cost-optimal**

cost-optimal algorithm

- how about adding the n/p numbers first locally by the p processors
 - this takes time $O(n/p)$
- now, we have p numbers to add, each number on 1 of the p processors
- we have shown before that this can be done on $O(\log p)$
- therefore, this resulting **running time** is $O(n/p + \log p)$
- the cost (running time multiplied by p) is
 - $O(n + p \log p)$
- if $n \in \Omega(p \log p)$, the cost is $O(n)$, which is the same as in the sequential case
 - thus, the algorithm is **cost-optimal**

- [1] A. Grama, G. Karypis, V. Kumar, and A. Gupta. *Introduction to Parallel Computing*. Second. Addison-Wesley, 2003. ISBN: 0201648652 9780201648652.